


# Network Support For Scalable And High-Performance Cloud Exchanges

Muhammad Haseeb, Jinkun Geng, Daniel Duclos-Cavalcanti, Xiyu Hao, Ulysses Butler, Radhika Mittal, Srinivas Narayana, Anirudh Sivaraman

New York University, Clockwork Inc., Technical University of Munich, Rutgers University, University of Illinois Urbana-Champaign





# Benefits of Public Cloud for Financial Exchanges

- Reducing constraints of physical space around the matching engine
- Unprecedented scale – 1000s of participants can be supported
- Low barriers to entry for launching new global markets
- Various analytics/ML services residing nearby
- And typical benefits of cloud: flexible resource allocation, reduced cost, offloading management etc.

**There will be trade-offs.**

- Latencies will be higher.
- Fairness guarantees will be coarser.


# Onyx: Scalable Cloud Financial Exchange

Multicast service that disseminates market data to 1000 participants, each receiving a message within 1-microsecond of each other


An order submission service which effectively handles bursty traffic while maintaining fairness of competition among the participants

Prototyped on AWS and GCP


# Primer On Financial Exchanges




# Primer On Financial Exchanges




# Primer On Financial Exchanges



# Primer On Financial Exchanges




# Primer On Financial Exchanges







# Primer On Financial Exchanges




# Primer On Financial Exchanges




# Primer On Financial Exchanges




# Primer On Financial Exchanges



# Primer On Financial Exchanges




# Primer On Financial Exchanges



# Primer On Financial Exchanges

Exchange Server




order closer to mid price → Immediate execution/match

order away from mid price → potential execution/match in future

Position in the queue/price level matters!

# Primer On Financial Exchanges

Exchange Server




Traditional exchanges built in private data centers or colocation facilities!

order arrival  
execution

Competition Among  
MPs

Position in the queue/price  
level matters!





## Public Cloud Exhibits High Latency Variance

- Latency varies temporally
- Latency varies spatially – across multiple clients
- Unfit for fair competition
- Onyx: mechanisms to achieve fairness + high performance at scale

# Fairness<sup>1</sup>

Outbound Fairness

Inbound Fairness

# Fairness<sup>1</sup>

Outbound Fairness




MP1 MP2 MP3

Three light blue ovals are arranged horizontally, each containing the text "MP1", "MP2", and "MP3" respectively.

Inbound Fairness

# Fairness<sup>1</sup>

Outbound Fairness



MP1 MP2 MP3

Three blue ovals are arranged horizontally at the bottom of the diagram, each containing the text "MP1", "MP2", and "MP3" respectively.

Inbound Fairness

# Fairness<sup>1</sup>

Outbound Fairness



Inbound Fairness

# Fairness<sup>1</sup>

## Outbound Fairness

Exchange  
Server

Difference in data  
reception time for any two  
participants should be 0

t1



MP1

t2



MP2

t3





MP3

## Inbound Fairness

# Fairness<sup>1</sup>



## Outbound Fairness



## Inbound Fairness

# Fairness<sup>1</sup>

## Outbound Fairness

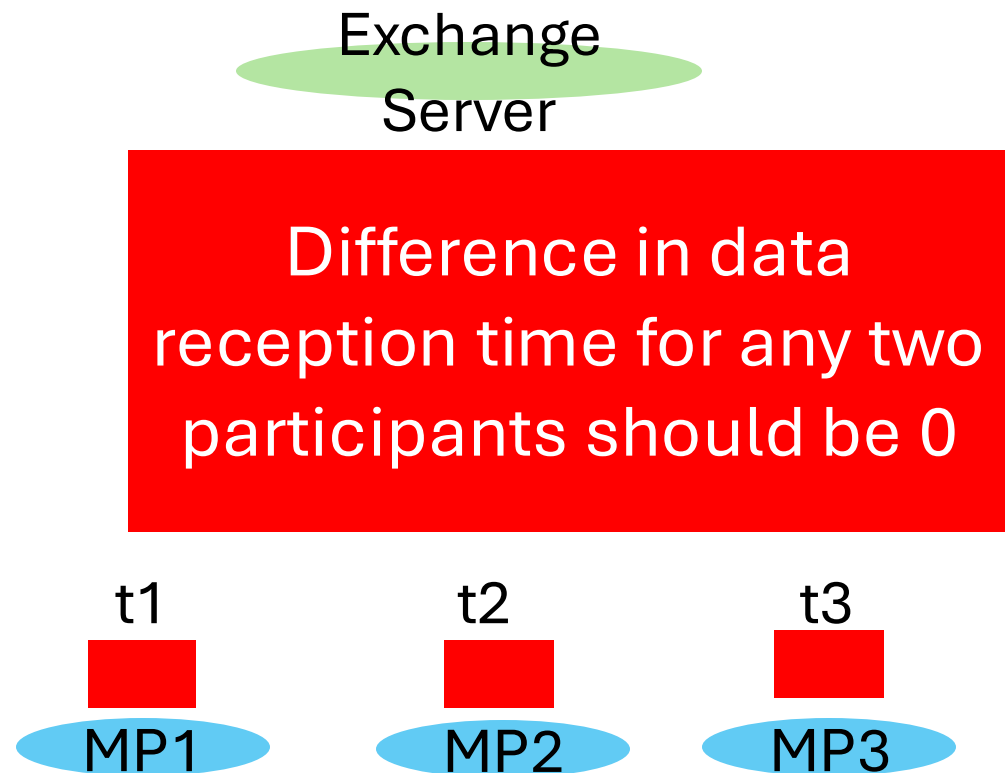


## Inbound Fairness



# Fairness<sup>1</sup>


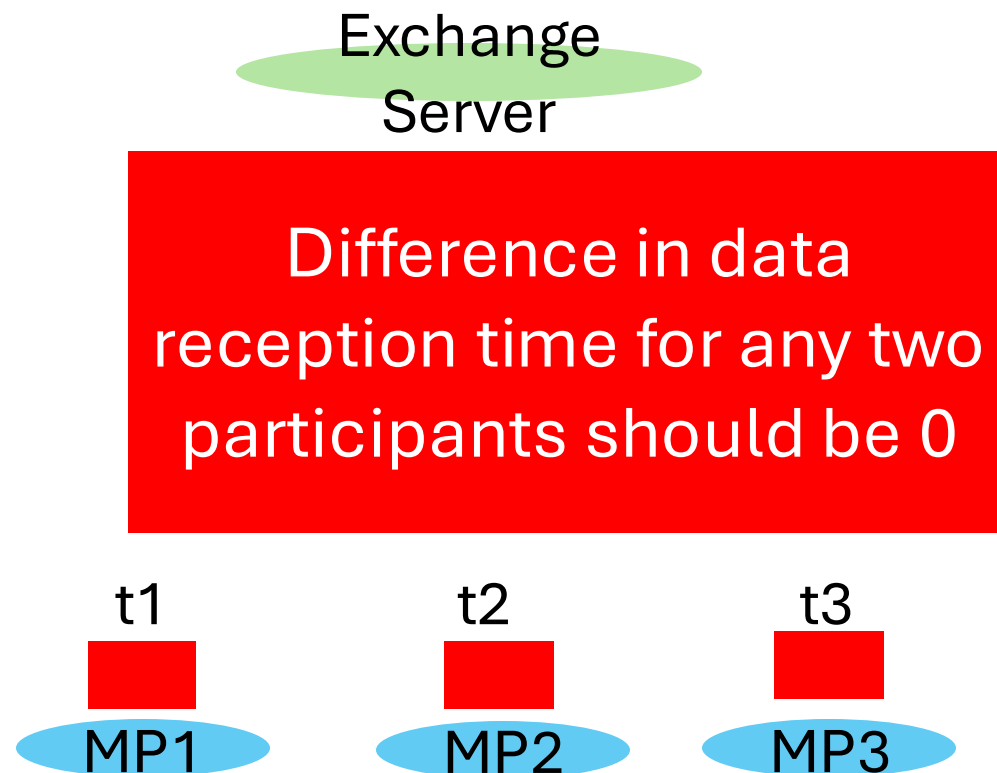
## Outbound Fairness



## Inbound Fairness

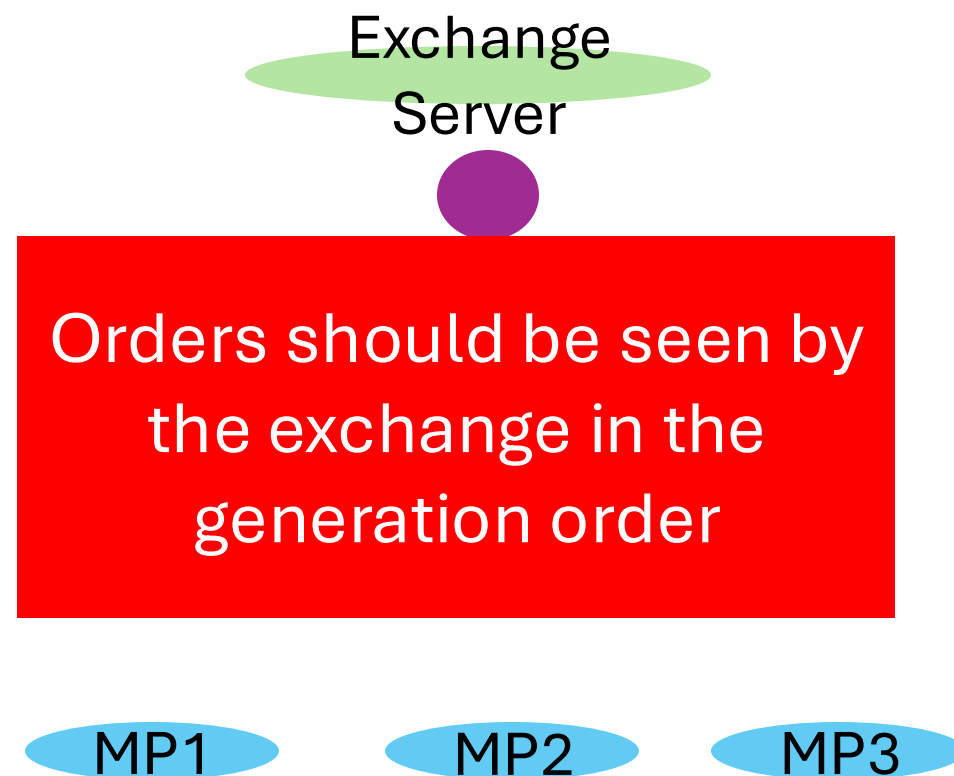
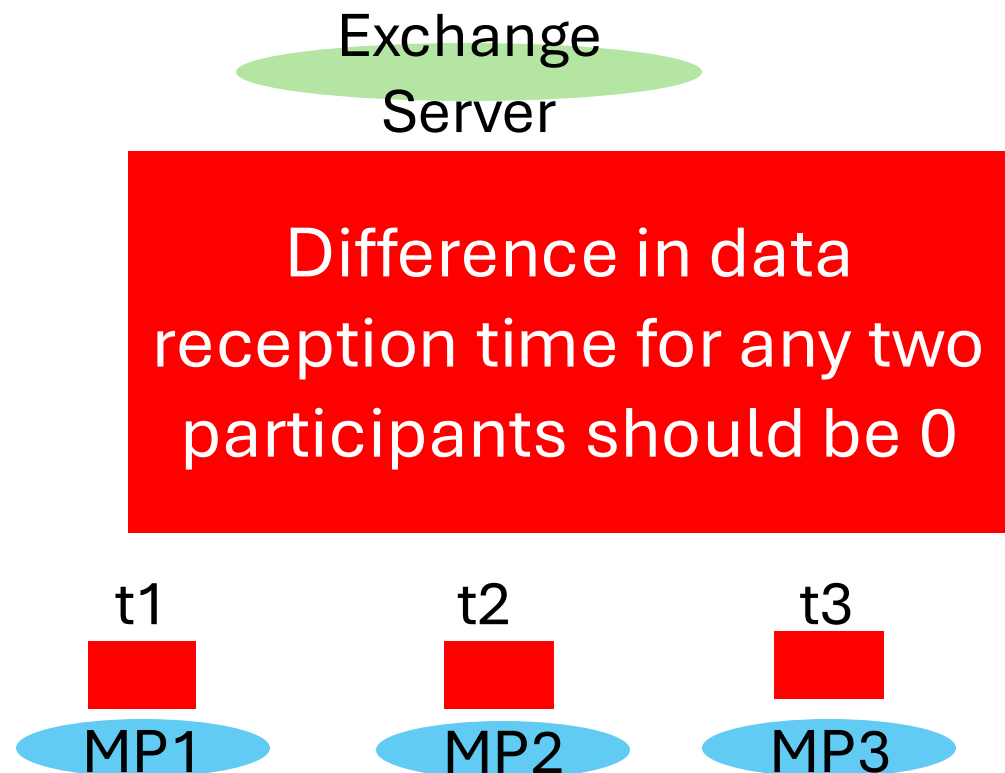
# Fairness<sup>1</sup>

## Outbound Fairness



# Fairness<sup>1</sup>

## Outbound Fairness



## Inbound Fairness

# Fairness<sup>1</sup>

## Outbound Fairness

Difference in data reception time for any two participants should be 0

Orders should be seen by the exchange in the generation order

## Inbound Fairness

# Fairness In Practice

## Outbound Fairness

Difference in data reception time for any two participants should be 0

Orders should be seen by the exchange in the generation order

## Inbound Fairness

# Fairness In Practice

## Outbound Fairness

Difference in data reception time for any two participants should be  $\sim 0$

Orders should be seen by the exchange in the generation order

## Inbound Fairness

# Fairness In Practice

## Outbound Fairness

Difference in data reception time for any two participants should be  $\sim 0$

Orders should be seen by the exchange in the timestamped generation order

## Inbound Fairness

Achieving fairness while being performant at scale is our goal.

# Fairness In Practice

## Outbound Fairness

Difference in data  
reception time for any two  
participants

We target 1000 market participants (MPs).  
Each MP is a separate VM.

Orders should be seen by  
the exchange in the  
timestamped generation  
order

## Inbound Fairness

Achieving fairness while being performant at  
**scale** is our goal.

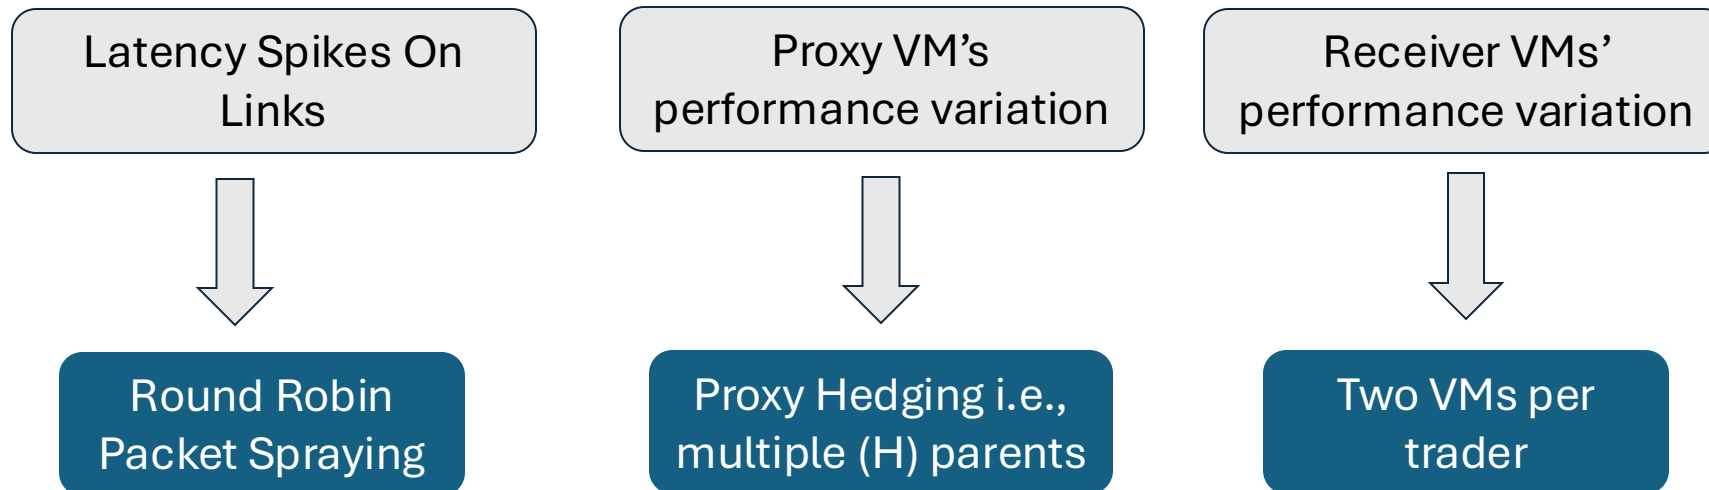


# Scalable Outbound Fairness

Achieved by a simple “hold-and-release” primitive.

We scale it much further by proposing an efficient market data multicast mechanism.

A tree of proxy nodes is utilized that helps scale multicast.



# Scalable Inbound Fairness



Sequencer

The diagram shows a dark blue rounded rectangle partially overlapping a light blue rounded rectangle. The word "Sequencer" is centered in the light blue area.




Scheduling  
Policy

The diagram shows a dark blue rounded rectangle partially overlapping a light blue rounded rectangle. The words "Scheduling" and "Policy" are centered in the light blue area, stacked vertically.


# A sequencer

- Guarantees that orders are seen by the exchange in non-decreasing order of their (generation) timestamps
- Assumes clock synchronization






# A sequencer

Clients attach message generation timestamps to outgoing messages



Exchange Server


Sequencer waits for at least one message from each client



Clients/Participants




# A sequencer

Clients attach message generation timestamps to outgoing messages



Exchange  
Server

Sequencer waits for at least one message from each client




Clients/Participants






# A sequencer

Clients attach message generation timestamps to outgoing messages



Exchange Server

Sequencer waits for at least one message from each client




Clients/Participants




# A sequencer

Clients attach message generation timestamps to outgoing messages






Exchange Server

Sequencer waits for at least one message from each client




It releases the smallest ts' message and waits again



Clients/Participants

# A sequencer

Clients attach message generation timestamps to outgoing messages






Exchange Server



Sequencer waits for at least one message from each client

It releases the smallest ts' message and waits again




Clients/Participants





# A sequencer

Clients attach message generation timestamps to outgoing messages



Exchange  
Server


Sequencer waits for at least one message from each client

It releases the smallest ts' message and waits again

Clients/Participants

# A sequencer

Clients attach message generation timestamps to outgoing messages



Sequencer waits for at least one message from each client

It releases the smallest ts' message and waits again

Clients/Participants

# A sequencer





Drawback: liveness (progress) may halt under client failures

Safety (fairness) is guaranteed if clocks are perfectly synchronized

In practice, we ensure that clock synchronization accuracy is sufficiently high compared to the time granularity of interest.


# Limit Order Queue

- During bursty market activity, the exchange server may be overwhelmed.
- LOQ schedules order in a way to avoid idling the matching engine, utilizing the knowledge that matching engine uses price time priority algorithm for matching orders




Naively doing such prioritization will not preserve inbound fairness.


# Limit Order Queue Gracefully Handles Bursts




# Limit Order Queue Preserves Fairness



# Limit Order Queue Preserves Fairness




# Limit Order Queue Preserves Fairness





# Limit Order Queue Preserves Fairness



# Limit Order Queue Preserves Fairness

$S_{\text{FIFO}}$



Exchange


$I_{\text{FIFO}}$

FIFO Fabric



Exchange

LOQ Fabric



# Limit Order Queue Preserves Fairness

$S_{\text{FIFO}}$



Exchange

$I_{\text{FIFO}}$

FIFO Fabric





Exchange


LOQ Fabric




# Limit Order Queue Preserves Fairness



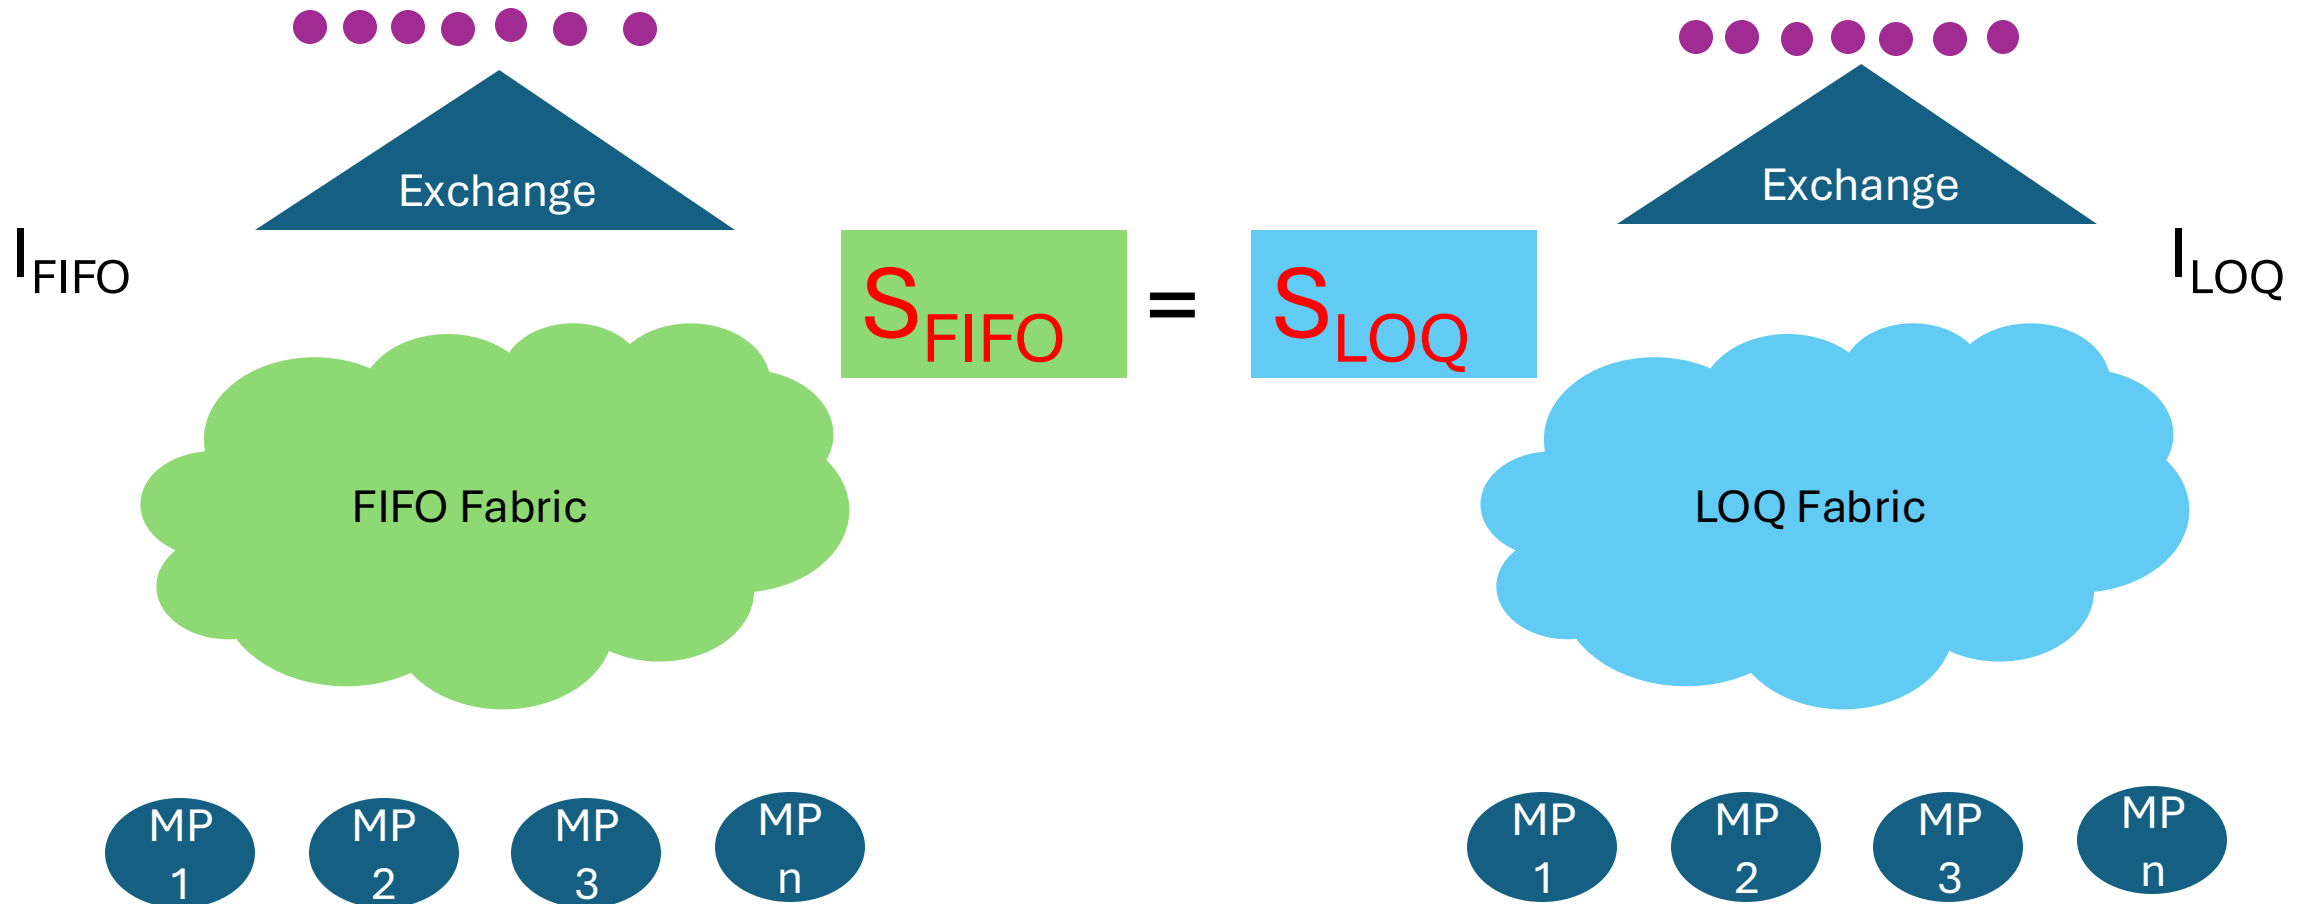
# Limit Order Queue Preserves Fairness




# Limit Order Queue Preserves Fairness



# Limit Order Queue Preserves Fairness



# Limit Order Queue Preserves Fairness

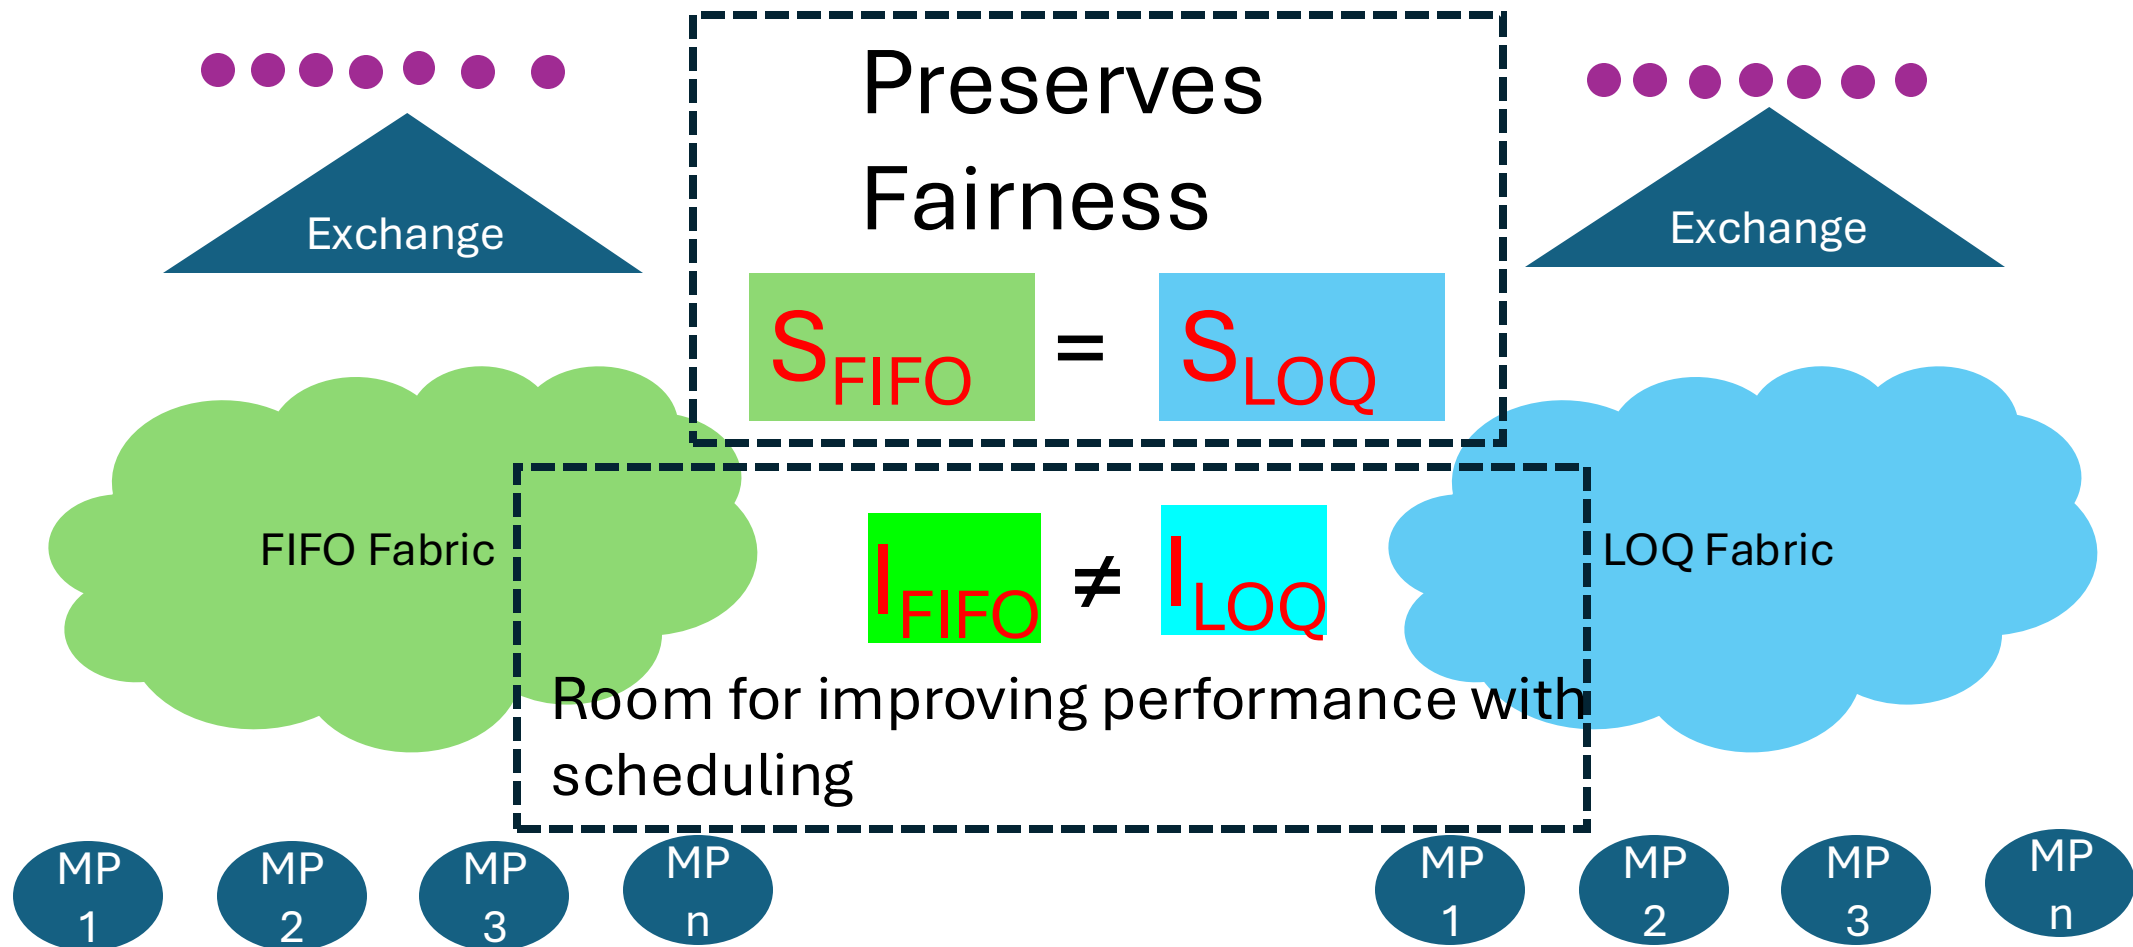




# Limit Order Queue Preserves Fairness




# Limit Order Queue Preserves Fairness






Evaluation


# Multicast Service



# Multicast Service



# Order Submission Service



# Concluding Remarks

- Cloud financial exchanges can be realized by developing new primitives, without specialized infrastructure
  - Some fairness/performance guarantees need to be accordingly relaxed
- Such exchanges present a new operating point in the cost-performance curve

